

# Performance prediction on heterogeneous systems using statistical methods

Spyrou Michalis  
University of Thessaly

# Purpose of this thesis

- Predicting performance on heterogeneous systems
- Creating the prediction model
- Facilitating the mapping of applications on various computational resources

# Our Approach

- Collection of hardware events
  - PAPI
  - Nvprof
- Use of Statistical methods
  - Linear Regression
  - Neural Networks
  - Random Forests


# Background Issues

# Architectural differences

- CPUs:
  - few processors supporting ~1 hardware thread
  - execute one stream of instructions as fast as possible
  - latency : large caches and branch prediction hardware
  - a few hundreds GFLOPS
- GPUs:
  - many processors supporting many hardware threads
  - execute many parallel streams of instructions as fast as possible
  - latency : supporting thousands of threads at once
  - thousands of GFLOPS

# Linear Regression

- Linear relationship
- Equation form:  $\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\varepsilon}$




# Neural Networks

- Interconnected **nodes**, forming **layers**

- Each node performs a mathematical function, such as

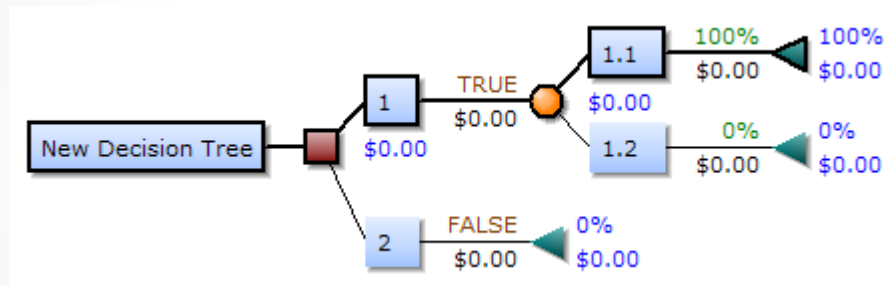
$$S(t) = \frac{1}{(1+e^{-t})}$$

- **Black box**



# Random Forests

- A combination of **decision trees**



- Node → **Test** on an attribute
- Branch → **Result** of a test
- Leaf → **Decision** taken



# Data Collection

# Benchmarks

- **Rodinia** benchmark suite 2.4
- The Scalable Heterogeneous Computing (**SHOC**) Benchmark Suite

# Intel Xeon CPUs

- Use of **OpenCL**
- Get number of compute units
- Use of low level **PAPI** Interface for events
- OpenCL built in functions for time

# Event measurement example

```
clFinish(command_queue);
```

```
start_cpu_counter ( counter, &EventSet,  
                   process_threads, compute_units);
```

```
error = clEnqueueNDRangeKernel( command_queue, kernel, 1,  
                                NULL, global_work_size, local_work_size, 0, NULL, NULL);
```

```
clFinish(command_queue);
```

```
stop_cpu_counter ( EventSet, &results, compute_units);
```

# NVIDIA GPUs

- **NVprof:** not supporting OpenCL
- Used CUDA
- CL Interface

# Model Formation

- Find **correlation** between target value and measured variables

Intel to NVIDIA		NVIDIA to Intel	
Counter	Correlation	Counter	Correlation
BR_TKN	0,8874	ldst_executed	0,88
TOT_CYC	0,868	inst_issued	0,88
BR_PRC	0,862	issue_slots	0,88
TOT_INS	0,861	inst_executed	0,86
LD_INS	0,8553	ldst_issued	0,86
BR_CN	0,8543	cf_issued	0,85
LST_INS	0,8503	cf_executed	0,85
BR_INS	0,8473	l2_read_trans	0,83

# Predicting OpenCL CUDA execution time

- Linear regression
- Tested different number of hardware events.  
Calculated **RMSE**

Number of Events Used	RMSE
9	5,12
8	7,29
7	7,17
6	8,86
5	229

# Predicting OpenCL CUDA execution time

- Linear regression formula:

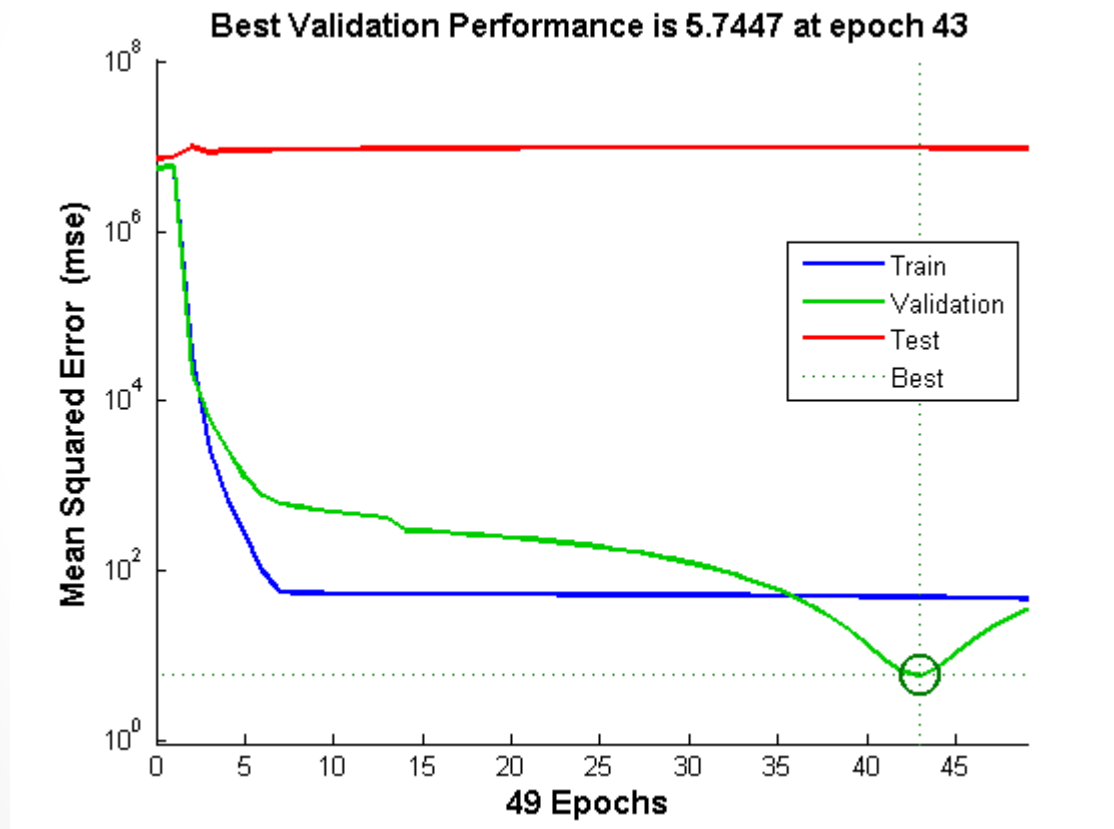
*Execution time* =

$$\begin{bmatrix} -2.9 \times 10^{-2} & -1.24 \times 10^{-3} & -1.12 & 7.865 \times 10^{-3} & -4.81 \times 10^{-3} & 1.08 \end{bmatrix} \times \begin{bmatrix} BR\_TKN \\ TOT\_CYC \\ BR\_PRC \\ TOT\_INS \\ LD\_INS \\ BR\_CN \end{bmatrix} \times 10^{-6} + 2.0829$$



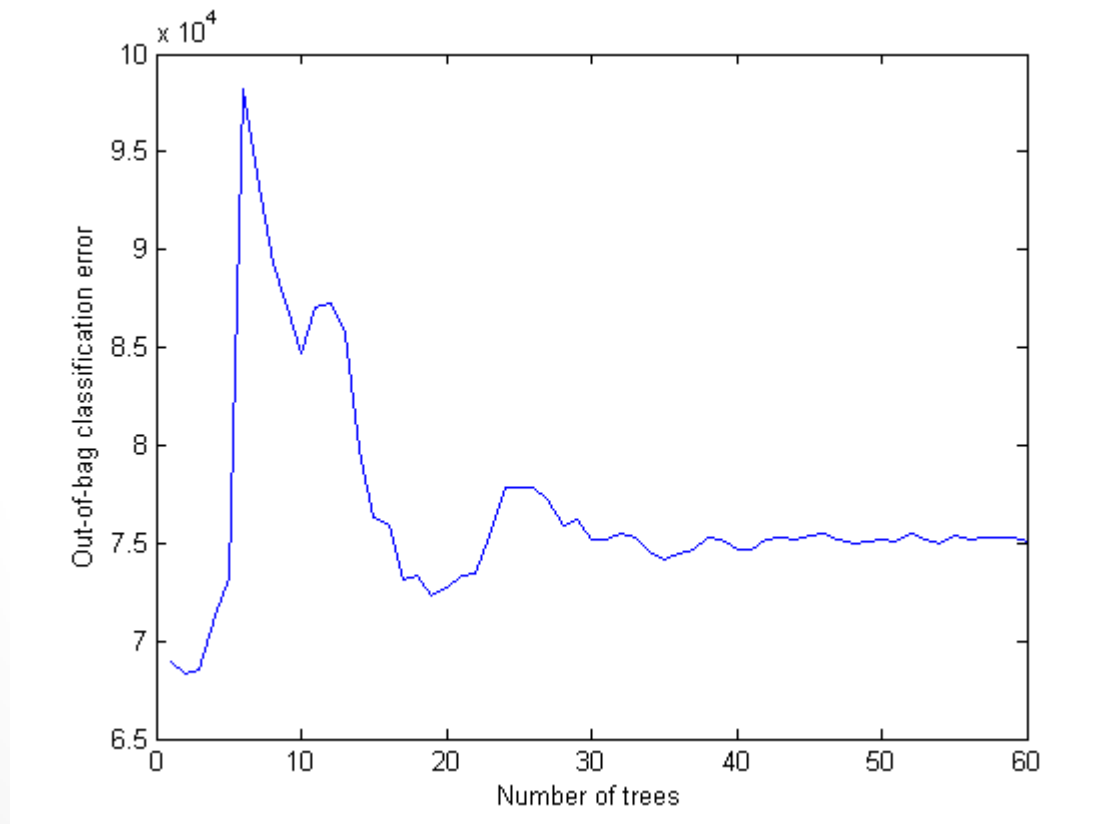
# Predicting OpenCL CUDA execution time

- Neural Networks



# Predicting OpenCL CUDA execution time

- Random forests



# Predicting Intel OpenCL execution time

- Linear regression

Number of Events Used	RMSE
10	92,6
9	99,1
8	110
7	110
6	109
5	115

# Predicting Intel OpenCL execution time


- Linear regression formula:

*Execution time* =

$$\begin{bmatrix} -2.12 & 65.06 & -65.9 & 1.07 & 0.4 & 7.96 & -7.89 & 0.37 & -35.8 & 43.4 \end{bmatrix} \times \begin{bmatrix} ldst\_executed \\ inst\_issued \\ issue\_slots \\ inst\_executed \\ ldst\_issued \\ cf\_issued \\ cf\_executed \\ l2\_read\_trans \\ l2\_write\_trans \\ dram\_write\_trans \end{bmatrix} \times 10^{-5} + 14.791$$


# Predicting Intel OpenCL execution time

- Neural Networks




# Predicting Intel OpenCL execution time

- Random Forests




# Evaluation

# Intel Xeon OpenCL to GTX480 CUDA - Linear

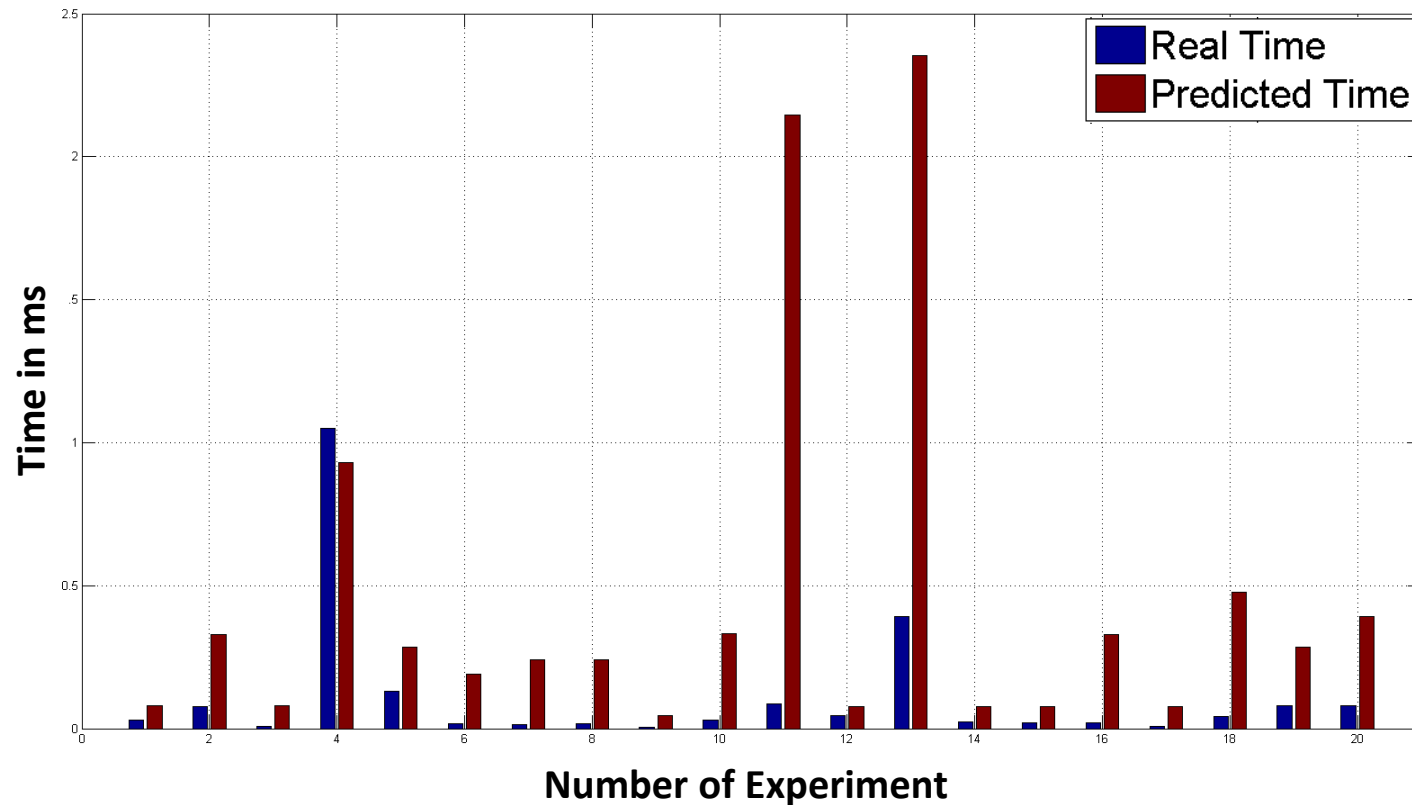





# Intel Xeon OpenCL to GTX480 CUDA - NN




# Intel Xeon OpenCL to GTX480 CUDA - RF




# GTX480 CUDA to Intel Xeon OpenCL - Linear



# GTX480 CUDA to Intel Xeon OpenCL - NN



# GTX480 CUDA to Intel Xeon OpenCL - RF



# Future work

- Predict execution time on ATI GPUs
- Increase size of training set
- Predict power
- Use our models with a run-time system

# Questions